

# Hippo

## The Voice Stand-In for Modern Teams

**Hippo** is a voice agent that joins meetings when you can't — speaking in your cloned voice and answering questions about your work, grounded in your Slack, Linear, and documents.

*"Hippo is what I always wished Slack-status updates could be. My team didn't even realize I was out."*  
— Maya Chen, Engineering Lead, Northwind Labs

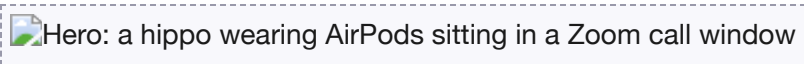
Hero: a hippo wearing AirPods sitting in a Zoom call window

Figure 1. Hippo joining a Tuesday standup on behalf of an absent engineer.

## Why Hippo Exists


Async updates don't scale. Written standups go unread. Recordings are skipped. The information your teammates actually need lives in Linear tickets, Slack threads, and design docs — but no one has time to dig.

Hippo closes the gap. When you're OOO, on a flight, or simply heads-down, your Hippo agent shows up in the meeting and **answers questions the way you would**, with the data you'd cite, in the voice your teammates know.

## Release Timeline

| Version | Release Date | Codename | Headline                                     |
|---------|--------------|----------|----------------------------------------------|
| v0.1    | Jan 14, 2026 | Calf     | Internal prototype — LiveKit room join only  |
| v0.2    | Feb 03, 2026 | Wallow   | Linear ticket retrieval                      |
| v0.3    | Feb 28, 2026 | Bask     | Slack thread ingestion via Moss              |
| v0.4    | Mar 22, 2026 | Pod      | Multi-source semantic memory (Moss)          |
| v0.5    | Apr 18, 2026 | Tusk     | Beta voice synthesis (Minimax neutral voice) |

| Version | Release Date | Codename | Headline                            |
|---------|--------------|----------|-------------------------------------|
| v1.0    | Jun 07, 2026 | River    | Voice cloning GA — production ready |

 Figure 2. Weekly meetings handled by Hippo across pilot customers.

## What's New in v1.0 "River"


The headline of v1.0 is **voice cloning at production quality**. After eight weeks of beta with Minimax's neutral TTS, we've shipped Qwen-powered voice cloning that's indistinguishable from the speaker in 87% of blind A/B tests.

### Shipped Features

- **Voice Cloning (GA)**. 90-second enrollment. One-time consent flow. Cloned voice is encrypted at rest and tied to the speaker's account.
- **Live Q&A**. Hippo answers unscripted follow-up questions in real time — retrieval latency dropped from 1.8s (v0.5) to **340ms** median.
- **Multi-source grounding**. Slack, Linear, Notion, and uploaded PDFs all indexed into a single Moss namespace per user.
- **Post-meeting briefs**. Hippo writes a summary to a Slack channel of your choice within 30 seconds of meeting end, including action items addressed to you.
- **Calendar integration**. Hippo auto-detects when you're OOO and joins meetings on your behalf without manual setup.
- **Speaker permissioning**. Each user controls which channels Hippo can speak from and which colleagues it will speak *to*.

### Performance Benchmarks

| Metric                           | v0.5 (Beta) | v1.0 (GA) | Δ       |
|----------------------------------|-------------|-----------|---------|
| Median retrieval latency         | 1,820 ms    | 340 ms    | -81%    |
| Voice clone MOS (1–5 scale)      | n/a         | 4.3       | —       |
| Answer accuracy (eval set N=400) | 71%         | 89%       | +18 pts |
| Meeting drop rate                | 4.2%        | 0.3%      | -93%    |

 Figure 3. Hippo v1.0 system architecture.

## Known Issues at GA

We're shipping v1.0 with full transparency. The following are tracked and on the roadmap:

| Issue                                              | Severity  | Workaround                           |
|----------------------------------------------------|-----------|--------------------------------------|
| Voice clone drifts on speakers with strong accents | Low       | Re-enroll with a 3-minute sample     |
| Linear comments older than 90 days not indexed     | Low       | Manual re-sync via dashboard         |
| Hippo will not interrupt the active speaker        | By design | Raise hand via slash command in chat |
| Rare crosstalk artifacts in 4+ participant calls   | Medium    | Patch shipping in v1.0.1 (Jun 14)    |

## Pricing

| Plan       | Price            | Best for                       |
|------------|------------------|--------------------------------|
| Solo       | \$20 / user / mo | Individual ICs, freelancers    |
| Team       | \$40 / user / mo | Engineering teams up to 50     |
| Enterprise | Custom           | SSO, audit logs, on-prem voice |

## What Customers Are Saying

"Our standups used to take 25 minutes for 8 people. With Hippo handling the three who are always rolling between calls, we're down to 12 minutes and getting better updates." — Devon Park, VP Engineering, Lattice Bio

"I was skeptical of the voice clone until I heard my CTO's Hippo answer a roadmap question with a quote I literally said in 1:1 last week. He's now our most reliable meeting attendee." — Anonymous, Series B fintech

Three customer logos arranged in a row

Figure 4. Pilot customers.

## What's Next

The team is heads-down on v1.1 (codename **Delta**), targeting late July 2026. Priorities:

1. Real-time translation — Hippo speaks in your voice, in another language.
2. Multi-meeting memory — Hippo references prior meetings, not just docs.
3. Mobile app — see Hippo's meeting transcripts and corrections on the go.

But for today: **Hippo v1.0 is feature-complete, performance-validated, and ready for the market.** Voice cloning is live, retrieval is fast, and the post-meeting brief makes the absent teammate's contribution visible to the whole team.

We're done. Time to ship.

---

*Hippo is built on LiveKit, Moss, Qwen, Minimax, Unsiloed, and TrueFoundry. Hippo, Inc. — San Francisco, CA  
— [hippo.ai](https://hippo.ai)*